

Analysis and Performance evaluation of Terapixel rendering in (Super)Cloud Computing Data

1.Introduction

The purpose of this study is to analyze the IoT environmental data collected by Newcastle Urban Observatory for the city of Newcastle Upon Tyne. The main objective is to evaluate and explore performance timings of render application and GPU card and in each task, the details of which part of the image was being rendered. Terapixel images are rendered using a scalable cloud-based visualization architecture. The Terapixel image, once created, allows for interactive exploration of the city and its data at a wide range of sensing scales ranging from the entire city to a separate desk in a room, and is accessible via a broad range of thin client devices. CRISP-DM (Cross- Industry Standard Process for Data Mining) model will be used in this data analysis. This project will be entirely dedicated to the EDA (Exploratory Data Analysis) process.

2. Data Exploration Planning and Analysis Requirement:

Tera scope terapixel data is subjected to preliminary analysis to better understand the data and provide information to business stakeholders. Based on the data set, this analysis of GPU cards and XY coordinates will aid in rendering Terapixel images in an efficient and effective manner.

2.1 Data Exploration Planning:

- Assessing the event types that dominate task runtimes.
- Examining the correlation between GPU temperature and performance.
- Analyzing the connection between increased power draw and render time.
- Identifying GPU cards (based on their serial numbers) whose performance differs to other cards.
- Exploring the effectiveness of the task scheduling process.

2.2 Analysis Plan and Requirement:

The analysis strategy for this report is to investigate the three data sets generated while different virtual machines render 3D images on 1024 GPU nodes during a run. This run is divided into three jobs to render the data visualization output, which show performance timing of the render application, performance of the GPU card, and details of which part of the image was rendered in each task. The requisite is to understand the data, then clean and preprocess the data before performing exploratory data analysis. This analysis will aid in the betterment of the rendering process.

Installing pandasql

```
!pip install pandasql
```

```
Requirement already satisfied: pandasql in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (0.7.3)
Requirement already satisfied: pandas in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from
pandasql) (1.4.4)
Requirement already satisfied: numpy in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from
pandasql) (1.21.5)
Requirement already satisfied: sqlalchemy in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from
pandasql) (1.4.39)
Requirement already satisfied: python-dateutil>=2.8.1 in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from pandas-
>pandasql) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from pandas-
>pandasql) (2022.1)
Requirement already satisfied: greenlet!=0.4.17 in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from
sqlalchemy->pandasql) (1.1.1)
Requirement already satisfied: six>=1.5 in
/Users/gowtham/opt/anaconda3/lib/python3.9/site-packages (from python-
dateutil>=2.8.1->pandas->pandasql) (1.16.0)
```

```
# Importing all the required libraries
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
import datetime
from pandasql import sqldf
from datetime import datetime
from pandas import Series
sqlfn = lambda q: sqldf(q, globals())
```

3. Data Understanding:

This analysis utilizes data from a dataset generated during a run with 1024 nodes. This analysis will be performed using three csv files application-checkpoints.csv, gpu.csv, task-x-y.csv. Below all the three csv files are read:

```
# Reading all the three CSV files
```

```
appcheckpoint=pd.read_csv("application-checkpoints.csv")
gpu=pd.read_csv("gpu.csv")
xy=pd.read_csv("task-x-y.csv")
```

Table **application-checkpoints.csv**:

Displaying first three rows of appcheckpoint table

```
appcheckpoint.head(3)
```

	timestamp	hostname	\
0	2018-11-08T07:41:55.921Z	0d56a730076643d585f77e00d2d8521a00000N	
1	2018-11-08T07:42:29.842Z	0d56a730076643d585f77e00d2d8521a00000N	
2	2018-11-08T07:42:29.845Z	0d56a730076643d585f77e00d2d8521a00000N	

	eventName	eventType	jobId	\
0	Tiling	STOP	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
1	Saving Config	START	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
2	Saving Config	STOP	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	

	taskId
0	b47f0263-ba1c-48a7-8d29-4bf021b72043
1	20fb9fcf-a927-4a4b-a64c-70258b66b42d
2	20fb9fcf-a927-4a4b-a64c-70258b66b42d

Displaying data types of columns in appcheckpoint table

```
appcheckpoint.dtypes
```

```
timestamp    object
hostname     object
eventName    object
eventType    object
jobId        object
taskId       object
dtype: object
```

Table **gpu.csv**

Displaying first three rows of GPU table

```
gpu.head(3)
```

	timestamp	hostname	\
0	2018-11-08T08:27:10.314Z	8b6a0eebc87b4cb2b0539e81075191b900001C	
1	2018-11-08T08:27:10.192Z	d8241877cd994572b46c861e5d144c85000000	
2	2018-11-08T08:27:10.842Z	db871cd77a544e13bc791a64a0c8ed50000006	

	gpuSerial	gpuUUID
0	323217055910	GPU-1d1602dc-f615-a7c7-ab53-fb4a7a479534

```

131.55
1 323617020295 GPU-04a2dea7-f4f1-12d0-b94d-996446746e6f
117.03
2 323217056562 GPU-f4597939-a0b4-e78a-2436-12dbab9a350f
121.64

```

```

      gpuTempC  gpuUtilPerc  gpuMemUtilPerc
0           48           92           53
1           40           92           48
2           45           91           44

```

Displaying data types of columns in gpu table

```
gpu.dtypes
```

```

timestamp      object
hostname        object
gpuSerial       int64
gpuUUID         object
powerDrawWatt   float64
gpuTempC        int64
gpuUtilPerc     int64
gpuMemUtilPerc  int64
dtype: object

```

Table **task-x-y.csv**

Displaying first three rows of task-x-y table

```
xy.head(3)
```

```

                                taskId \
0  00004e77-304c-4fbd-88a1-1346ef947567
1  0002afb5-d05e-4da9-bd53-7b6dc19ea6d4
2  0003c380-4db9-49fb-8e1c-6f8ae466ad85

                                jobId    x    y  level
0  1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705  116  178    12
1  1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705  142  190    12
2  1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705  142   86    12

```

Displaying data types of columns in xy table

```
xy.dtypes
```

```

taskId      object
jobId       object
x            int64
y            int64
level       int64
dtype: object

```

```
# Displaying the number of rows and columns in all the three tables
print("Number of rows and columns in appcheckpoint table :
",appcheckpoint.shape)
print("Number of rows and columns in gpu table : ",gpu.shape)
print("Number of rows and columns in xy table : ",xy.shape)
```

```
Number of rows and columns in appcheckpoint table : (660400, 6)
Number of rows and columns in gpu table : (1543681, 8)
Number of rows and columns in xy table : (65793, 5)
```

4. Data Preparation:

After Data Understanding, the data is cleaned and preprocessed in preparation for further data analysis. There are some redundant (duplicate) data values in this data set that must be cleaned before further processing.

Count of all the duplicates rows in all the three tables are displayed

```
# Count of duplicate records in appcheckpoint table
```

```
appcheckpoint.duplicated().sum()
```

```
2470
```

```
# Count of duplicate records in gpu table
```

```
gpu.duplicated().sum()
```

```
9
```

```
# Count of duplicate records in xy table
```

```
xy.duplicated().sum()
```

```
0
```

Duplicate records from appcheckpoint table is dropped

```
# Removing duplicate records from appcheckpoint table
```

```
appcheck = appcheckpoint.drop_duplicates()
appcheck
```

```

                                timestamp
hostname \
0      2018-11-08T07:41:55.921Z
0d56a730076643d585f77e00d2d8521a00000N
1      2018-11-08T07:42:29.842Z
0d56a730076643d585f77e00d2d8521a00000N
2      2018-11-08T07:42:29.845Z
0d56a730076643d585f77e00d2d8521a00000N
3      2018-11-08T07:42:29.845Z
```

0d56a730076643d585f77e00d2d8521a00000N
4 2018-11-08T07:43:13.957Z
0d56a730076643d585f77e00d2d8521a00000N
...
...
660395 2018-11-08T08:30:10.296Z
0745914f4de046078517041d70b22fe700000I
660396 2018-11-08T08:30:10.325Z
0745914f4de046078517041d70b22fe700000I
660397 2018-11-08T08:30:14.081Z
a77ef58b13ad4c01b769dac8409af3f800000H
660398 2018-11-08T08:30:14.127Z
a77ef58b13ad4c01b769dac8409af3f800000H
660399 2018-11-08T08:30:12.159Z
0d56a730076643d585f77e00d2d8521a000009

	eventName	eventType	\
0	Tiling	STOP	
1	Saving Config	START	
2	Saving Config	STOP	
3	Render	START	
4	TotalRender	STOP	
...	
660395	Tiling	STOP	
660396	Uploading	STOP	
660397	Tiling	STOP	
660398	TotalRender	STOP	
660399	Uploading	STOP	

	jobId	\
0	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
1	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
2	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
3	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
4	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
...	...	
660395	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
660396	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
660397	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
660398	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
660399	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	

	taskId
0	b47f0263-ba1c-48a7-8d29-4bf021b72043
1	20fb9fcf-a927-4a4b-a64c-70258b66b42d
2	20fb9fcf-a927-4a4b-a64c-70258b66b42d
3	20fb9fcf-a927-4a4b-a64c-70258b66b42d
4	20fb9fcf-a927-4a4b-a64c-70258b66b42d
...	...
660395	8261c0ff-03d6-48b3-a50f-da41cb3291fd

```

660396 8261c0ff-03d6-48b3-a50f-da41cb3291fd
660397 ce97e3e9-494a-43a7-aa85-edd2db4cf099
660398 ce97e3e9-494a-43a7-aa85-edd2db4cf099
660399 bbf4ff73-c754-4111-95cd-a339c87a2b17

```

[657930 rows x 6 columns]

Duplicate records from gpu table is dropped

Removing duplicate records from gpu table

```

gpufinal = gpu.drop_duplicates()
gpufinal

```

```

                                timestamp
hostname \
0      2018-11-08T08:27:10.314Z
8b6a0eebc87b4cb2b0539e81075191b900001C
1      2018-11-08T08:27:10.192Z
d8241877cd994572b46c861e5d144c85000000
2      2018-11-08T08:27:10.842Z
db871cd77a544e13bc791a64a0c8ed50000006
3      2018-11-08T08:27:10.424Z
b9a1fa7ae2f74eb68f25f607980f97d7000010
4      2018-11-08T08:27:10.937Z
db871cd77a544e13bc791a64a0c8ed50000003
...
...
1543676 2018-11-08T08:31:18.873Z
0d56a730076643d585f77e00d2d8521a00000N
1543677 2018-11-08T08:31:24.933Z
0d56a730076643d585f77e00d2d8521a00000N
1543678 2018-11-08T08:31:32.998Z
0d56a730076643d585f77e00d2d8521a00000N
1543679 2018-11-08T08:31:39.057Z
0d56a730076643d585f77e00d2d8521a00000N
1543680 2018-11-08T08:31:45.108Z
0d56a730076643d585f77e00d2d8521a00000N

```

```

                                gpuSerial                                gpuUUID \
0      323217055910  GPU-1d1602dc-f615-a7c7-ab53-fb4a7a479534
1      323617020295  GPU-04a2dea7-f4f1-12d0-b94d-996446746e6f
2      323217056562  GPU-f4597939-a0b4-e78a-2436-12dbab9a350f
3      325217085931  GPU-ad773c69-c386-a4be-b214-1ea4fc6045df
4      323217056464  GPU-2d4eed64-4ca8-f12c-24bc-28f036493ea2
...
...
1543676 325217086221  GPU-1265fef9-aea4-4a5e-8a63-cc5af7b19f4f
1543677 325217086221  GPU-1265fef9-aea4-4a5e-8a63-cc5af7b19f4f
1543678 325217086221  GPU-1265fef9-aea4-4a5e-8a63-cc5af7b19f4f
1543679 325217086221  GPU-1265fef9-aea4-4a5e-8a63-cc5af7b19f4f

```

```
1543680 325217086221 GPU-1265fef9-aea4-4a5e-8a63-cc5af7b19f4f
```

	powerDrawWatt	gpuTempC	gpuUtilPerc	gpuMemUtilPerc
0	131.55	48	92	53
1	117.03	40	92	48
2	121.64	45	91	44
3	50.23	38	90	43
4	141.82	41	90	47
...
1543676	26.99	30	0	0
1543677	26.90	29	0	0
1543678	26.88	29	0	0
1543679	26.89	29	0	0
1543680	26.89	29	0	0

```
[1543672 rows x 8 columns]
```

```
Duplicate records from xy table is dropped
```

```
# Removing duplicate records from xy table
```

```
xyfinal = xy.drop_duplicates()
```

```
xyfinal
```

	taskId \
0	00004e77-304c-4fbd-88a1-1346ef947567
1	0002afb5-d05e-4da9-bd53-7b6dc19ea6d4
2	0003c380-4db9-49fb-8e1c-6f8ae466ad85
3	000993b6-fc88-489d-a4ca-0a44fd800bd3
4	000b158b-0ba3-4dca-bf5b-1b3bd5c28207
...	...
65788	fce56316-25a6-44b8-b0cc-d135bc84deea
65789	feb48593-c70e-49b0-b19c-a4e2a5d2d760
65790	feb580b7-e879-40f2-aa57-28f3872af561
65791	ff1758b0-0a39-4f72-be8e-8ead79d691a6
65792	ff17a467-9962-4f3d-8e5e-5957efc3e8ab

	level	jobId	x	y
0	12	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	116	178
1	12	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	142	190
2	12	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	142	86
3	12	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	235	11
4	12	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	171	53
...


```

.
65788  1024-lvl8-5ad819e1-fbf2-42e0-8f16-a3baca825a63  15    2
8
65789  1024-lvl8-5ad819e1-fbf2-42e0-8f16-a3baca825a63    5   14
8
65790  1024-lvl8-5ad819e1-fbf2-42e0-8f16-a3baca825a63    3    6
8
65791  1024-lvl8-5ad819e1-fbf2-42e0-8f16-a3baca825a63    1    2
8
65792  1024-lvl8-5ad819e1-fbf2-42e0-8f16-a3baca825a63    7   12
8

```

```
[65793 rows x 5 columns]
```

```
# Change the datatype of a column
```

```
appcheck["timestamp"] = appcheck["timestamp"].apply(lambda x:
pd.Timestamp(x))
```

```
# Display the datatype of each column
```

```
appcheck.dtypes
```

```

timestamp      datetime64[ns, UTC]
hostname                object
eventName         object
eventType         object
jobId              object
taskId            object
dtype: object

```

```
# Change the datatype
```

```
gpufinal["timestamp"] = gpufinal["timestamp"].apply(lambda x:
pd.Timestamp(x))
```

```
# Change the datatype
```

```
gpufinal["gpuSerial"] = gpufinal["gpuSerial"].apply(lambda x: str(x))
```

```
# Display the datatype of each column
```

```
gpufinal.dtypes
```

```

timestamp      datetime64[ns, UTC]
hostname                object
gpuSerial         object
gpuUUID          object
powerDrawWatt      float64
gpuTempC           int64
gpuUtilPerc        int64
gpuMemUtilPerc     int64
dtype: object

```

```
# Create a new dataframe by filtering
```

```
appcheck_start = appcheck.loc[appcheck["eventType"] == "START"]
```

```

# Create a new dataframe by filtering
appcheck_stop = appcheck.loc[appcheck["eventType"] == "STOP"]

# Merge(Join) dataframes on multiple columns
appcheck_start_stop = pd.merge(appcheck_start, appcheck_stop,
                                how='left', on=['hostname', 'eventName', 'jobId', 'taskId'],
                                suffixes=('_Start_Time', '_Stop_Time'))

# # Display the first 3 rows
appcheck_start_stop.tail(4)

```

```

                                timestamp_Start_Time \
328961 2018-11-08 08:30:11.174000+00:00
328962 2018-11-08 08:30:08.759000+00:00
328963 2018-11-08 08:30:09.642000+00:00
328964 2018-11-08 08:30:09.649000+00:00

```

```

                                hostname  eventName \
328961  b9a1fa7ae2f74eb68f25f607980f97d7000005  Uploading
328962  6139a35676de44d6b61ec247f0ed865700001D    Tiling
328963  0d56a730076643d585f77e00d2d8521a00000D  Uploading
328964  0d56a730076643d585f77e00d2d8521a00000D    Tiling

```

```

                                eventType_Start_Time
jobId \
328961          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
328962          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
328963          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
328964          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705

```

```

                                taskId
timestamp_Stop_Time \
328961  120e5f6f-67a7-4973-b431-026f1a68400c 2018-11-08
08:30:12.074000+00:00
328962  30caa975-249a-445d-a6c0-61fa38d4c4bf 2018-11-08
08:30:09.730000+00:00
328963  423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6 2018-11-08
08:30:10.614000+00:00
328964  423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6 2018-11-08
08:30:10.548000+00:00

```

```

                                eventType_Stop_Time
328961          STOP
328962          STOP
328963          STOP
328964          STOP

```

```
# Create a new column from other two columns
appcheck_start_stop["Event_RenderTime"] =
(appcheck_start_stop["timestamp_Stop_Time"] -
appcheck_start_stop["timestamp_Start_Time"]).dt.total_seconds()
```

```
# Display datatype of each column
appcheck_start_stop.dtypes
```

```
timestamp_Start_Time    datetime64[ns, UTC]
hostname                object
eventName               object
eventType_Start_Time    object
jobId                   object
taskId                  object
timestamp_Stop_Time     datetime64[ns, UTC]
eventType_Stop_Time     object
Event_RenderTime        float64
dtype: object
```

```
# Display first 2 rows
```

```
appcheck_start_stop.tail(3)
```

```
timestamp_Start_Time \
328962 2018-11-08 08:30:08.759000+00:00
328963 2018-11-08 08:30:09.642000+00:00
328964 2018-11-08 08:30:09.649000+00:00
```

```
hostname  eventName \
328962    6139a35676de44d6b61ec247f0ed865700001D    Tiling
328963    0d56a730076643d585f77e00d2d8521a00000D    Uploading
328964    0d56a730076643d585f77e00d2d8521a00000D    Tiling
```

```
eventType_Start_Time
jobId \
328962          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
328963          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
328964          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
```

```
taskId
timestamp_Stop_Time \
328962  30caa975-249a-445d-a6c0-61fa38d4c4bf 2018-11-08
08:30:09.730000+00:00
328963  423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6 2018-11-08
08:30:10.614000+00:00
328964  423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6 2018-11-08
08:30:10.548000+00:00
```

	eventType_Stop_Time	Event_RenderTime
328962	STOP	0.971
328963	STOP	0.972
328964	STOP	0.899

Change the datatype

```
appcheck_start_stop["Event_RenderTime"] =
appcheck_start_stop["Event_RenderTime"].astype(str)
```

Change the datatype

```
appcheck_start_stop["Event_RenderTime"] =
appcheck_start_stop["Event_RenderTime"].astype('float32')
```

Display first 2 rows

```
appcheck_start_stop.tail(3)
```

	timestamp_Start_Time	\
328962	2018-11-08 08:30:08.759000+00:00	
328963	2018-11-08 08:30:09.642000+00:00	
328964	2018-11-08 08:30:09.649000+00:00	

	hostname	eventName	\
328962	6139a35676de44d6b61ec247f0ed865700001D	Tiling	
328963	0d56a730076643d585f77e00d2d8521a00000D	Uploading	
328964	0d56a730076643d585f77e00d2d8521a00000D	Tiling	

	eventType_Start_Time	jobId	\
328962	START	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
328963	START	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	
328964	START	1024-lvl12-7e026be3-5fd0-48ee-b7d1-abd61f747705	

	taskId	timestamp_Stop_Time	\
328962	30caa975-249a-445d-a6c0-61fa38d4c4bf	2018-11-08 08:30:09.730000+00:00	
328963	423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6	2018-11-08 08:30:10.614000+00:00	
328964	423b8511-cb2e-4aa4-bb5c-85ca4a2b7ac6	2018-11-08 08:30:10.548000+00:00	

	eventType_Stop_Time	Event_RenderTime
328962	STOP	0.971
328963	STOP	0.972
328964	STOP	0.899

```
# Sort the dataframe based on a column
appcheck_start_stop.sort_values(by=["hostname"], inplace=True)
```

```
# Display first 5 rows
appcheck_start_stop.tail(7)
```

```

                                timestamp_Start_Time \
285814 2018-11-08 07:46:50.935000+00:00
300124 2018-11-08 08:21:20.994000+00:00
218933 2018-11-08 07:43:33.460000+00:00
218938 2018-11-08 07:43:30.381000+00:00
277668 2018-11-08 08:00:45.782000+00:00
194757 2018-11-08 08:19:59.217000+00:00
229631 2018-11-08 08:13:37.033000+00:00
```

```

                                hostname  eventName \
285814 e7adc42d28814e518e9601ac2329c51300001D  Uploading
300124 e7adc42d28814e518e9601ac2329c51300001D    Render
218933 e7adc42d28814e518e9601ac2329c51300001D    Render
218938 e7adc42d28814e518e9601ac2329c51300001D  Uploading
277668 e7adc42d28814e518e9601ac2329c51300001D  Uploading
194757 e7adc42d28814e518e9601ac2329c51300001D    Render
229631 e7adc42d28814e518e9601ac2329c51300001D  Uploading
```

```

                                eventType_Start_Time
jobId \
285814          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
300124          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
218933          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
218938          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
277668          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
194757          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
229631          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
```

```

                                taskId
timestamp_Stop_Time \
285814 ad543e64-be07-4294-b562-c0588639d861 2018-11-08
07:46:52.016000+00:00
300124 ea7c0679-ff57-480f-a8b1-507c094b7c65 2018-11-08
08:21:59.023000+00:00
218933 de595141-d9c8-4e35-ac71-f14d1c9ebe80 2018-11-08
07:44:20.792000+00:00
218938 472afcae-6f17-43c1-8f1c-8df4116cfed7 2018-11-08
```

```

07:43:31.506000+00:00
277668 4001ba74-396f-4c68-80cb-165723760b6a 2018-11-08
08:00:46.780000+00:00
194757 4b2777ea-b877-498f-bace-43d867a86046 2018-11-08
08:20:33.780000+00:00
229631 69bfa379-618d-4d74-b9ad-ddb7d7765837 2018-11-08
08:13:38.079000+00:00

```

```

      eventType_Stop_Time  Event_RenderTime
285814                STOP              1.081000
300124                STOP              38.028999
218933                STOP              47.332001
218938                STOP              1.125000
277668                STOP              0.998000
194757                STOP              34.563000
229631                STOP              1.046000

```

Filter the dataframe

```

host_PT = appcheck_start_stop.loc[appcheck_start_stop["eventName"] ==
"TotalRender"]

```

Group the dataframe and average the other numerical columns

```

host_PT = host_PT.groupby(by=["hostname"], as_index=False).mean()

```

Display the last 3 rows

```

host_PT.head(3)

```

```

      hostname  Event_RenderTime
0  04dc4e9647154250beeee51b866b0715000000  44.350327
1  04dc4e9647154250beeee51b866b0715000001  45.698051
2  04dc4e9647154250beeee51b866b0715000002  40.275421

```

Display all column names

```

gpufinal.columns

```

```

Index(['timestamp', 'hostname', 'gpuSerial', 'gpuUUID',
      'powerDrawWatt',
      'gpuTempC', 'gpuUtilPerc', 'gpuMemUtilPerc'],
      dtype='object')

```

Group the dataframe and average the other numerical columns

```

gpu_PF =
gpufinal[["hostname", "gpuUUID", "gpuSerial", "powerDrawWatt", "gpuTempC",
"gpuUtilPerc", "gpuMemUtilPerc"]].groupby(by=["hostname", "gpuUUID", "gpu
Serial"], as_index=False).mean()

```

Display the first 3 rows

```

gpu_PF.tail(5)

```

```

      hostname \
1019  e7adc42d28814e518e9601ac2329c513000019

```

```

1020 e7adc42d28814e518e9601ac2329c51300001A
1021 e7adc42d28814e518e9601ac2329c51300001B
1022 e7adc42d28814e518e9601ac2329c51300001C
1023 e7adc42d28814e518e9601ac2329c51300001D

```

```

                                gpuUUID      gpuSerial
powerDrawWatt \
1019 GPU-53f2e984-36f3-a70f-0944-752868d6a833 320118119210
89.386562
1020 GPU-05e3f5b6-553b-1a37-bf7a-271c5b999a64 325117063265
82.943504
1021 GPU-8646ce3a-be99-02fa-a138-dafef7ea3cf2 325017048638
88.354997
1022 GPU-b0b09148-0762-390d-08d6-93c1a5ccd768 325117173230
88.411319
1023 GPU-0cee5a9f-749e-7780-791a-ff2b29590a38 320118119027
97.399313

```

```

      gpuTempC  gpuUtilPerc  gpuMemUtilPerc
1019 39.095270    64.512325    36.449034
1020 37.712858    61.762825    30.990007
1021 38.569046    63.919947    35.593062
1022 37.205863    64.930047    36.289141
1023 39.816000    64.793333    36.607333

```

Merge (Join) two dataframes

```
TP = pd.merge(host_PT, gpu_PF, on="hostname")
```

Display first 5 rows

```
TP.head(3)
```

```

                                hostname  Event_RenderTime \
0  04dc4e9647154250beeee51b866b0715000000    44.350327
1  04dc4e9647154250beeee51b866b0715000001    45.698051
2  04dc4e9647154250beeee51b866b0715000002    40.275421

```

```

                                gpuUUID      gpuSerial
powerDrawWatt \
0 GPU-a1119ee9-9cd1-919f-a479-b902142c717d 323217056165
95.868947
1 GPU-3dc1601f-0e52-2e31-6b8d-8537d356b84d 323617042956
91.813693
2 GPU-e2deaca4-3041-9bbd-b6ac-6d36e09ab116 323617021222
82.537798

```

```

      gpuTempC  gpuUtilPerc  gpuMemUtilPerc
0  43.525333    63.602667    35.876000
1  40.992000    64.630000    35.510000
2  38.048193    61.548862    30.505355

```

```
appcheck_start_stop.head(3)
```

```
          timestamp_Start_Time \
190313  2018-11-08 08:16:43.895000+00:00
283469  2018-11-08 07:47:22.860000+00:00
283468  2018-11-08 07:47:22.863000+00:00
```

```
          hostname          eventName \
190313  04dc4e9647154250beeee51b866b0715000000    Tiling
283469  04dc4e9647154250beeee51b866b0715000000  Saving Config
283468  04dc4e9647154250beeee51b866b0715000000    Render
```

```
          eventType_Start_Time
jobId \
190313          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
283469          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
283468          START  1024-lvl12-7e026be3-5fd0-48ee-b7d1-
abd61f747705
```

```
          taskId
timestamp_Stop_Time \
190313  a164499b-4bb8-42b7-9dc0-cf049b62731e 2018-11-08
08:16:44.897000+00:00
283469  ffd69831-0ba4-457b-b8a8-e37c49779d94 2018-11-08
07:47:22.863000+00:00
283468  ffd69831-0ba4-457b-b8a8-e37c49779d94 2018-11-08
07:47:55.554000+00:00
```

```
          eventType_Stop_Time  Event_RenderTime
190313          STOP          1.002000
283469          STOP          0.003000
283468          STOP          32.691002
```

```
# Merge two dataframes on multiple columns
```

```
xyfinal1 = pd.merge(xyfinal,
appcheck_start_stop.loc[appcheck_start_stop["eventName"]=="TotalRender
"],
                    how="inner", on=["jobId","taskId"],
suffixes=("_task","_ap"))
```

```
# Filter and create a new dataframe
```

```
xy_level8 =
xyfinal1[["Event_RenderTime","x","y"]].loc[xyfinal1["level"]==8]
```

```
# Group and aggregate the dataframe
```

```
xy_level8 = xy_level8.groupby(by=["x","y"], as_index=False).mean()
xy_level8.head(5)
```


	x	y	Event_RenderTime
0	0	0	43.853001
1	0	1	24.809999
2	0	2	27.111000
3	0	3	35.481998
4	0	4	43.898998

```
# Create a new dataframe using specific columns from other dataframe
Level_RT = xyfinal1[["level", "Event_RenderTime"]]
```

```
# Group and aggregate the dataframe
Level_RT = Level_RT.groupby(by=["level"], as_index=False).mean()
```

```
# Display the dataframe
Level_RT
```

	level	Event_RenderTime
0	4	52.181999
1	8	48.459820
2	12	42.581760

Data Analysis

1. Which tasks dominate even run times

```
# Group the dataframe
```

```
Event_Tab =
appcheck_start_stop[["eventName", "Event_RenderTime"]].groupby(by="eventName", as_index=False).mean()
```

```
# Sort the dataframe
```

```
Event_Tab.sort_values(by="Event_RenderTime", ascending=False,
inplace=True)
```

```
# Reset dataframe index to default
```

```
Event_Tab = Event_Tab.reset_index(drop=True)
```

```
# Display the dataframe
```

```
Event_Tab
```

	eventName	Event_RenderTime
0	TotalRender	42.604778
1	Render	41.208221
2	Uploading	1.393641
3	Tiling	0.973207
4	Saving Config	0.002476

```
Event_Tab = Event_Tab.sort_values(by='Event_RenderTime',
ascending=True)
```

```
# Plot a barplot using seaborn library
```

```
sns.set_style('whitegrid')
```


```

sns.set_context('notebook', font_scale=1.0)
plt.subplots(figsize=(6,5))
palette = {i: "red" if i == "TotalRender" else 'black' for i in
Event_Tab["eventName"]}
RenderTime_Plot = sns.barplot(data=Event_Tab, x="eventName",
y="Event_RenderTime",palette=palette)

# Add respective value on each bar
for g in RenderTime_Plot.patches:
    RenderTime_Plot.annotate(format(g.get_height(), '.7f'),
                             (g.get_x() + g.get_width() / 2.,
g.get_height()),
                             ha = 'center', va = 'center',
                             xytext = (0, 6),
                             textcoords = 'offset points')

# Change characteristics of the visual
plt.title("EventRenderTime", size=23)
plt.xlabel("Events", size = 16)
plt.ylabel("Average Render Time of Events(seconds)", size = 15)
plt.show()

```



2. Comparisons of GPU metrics

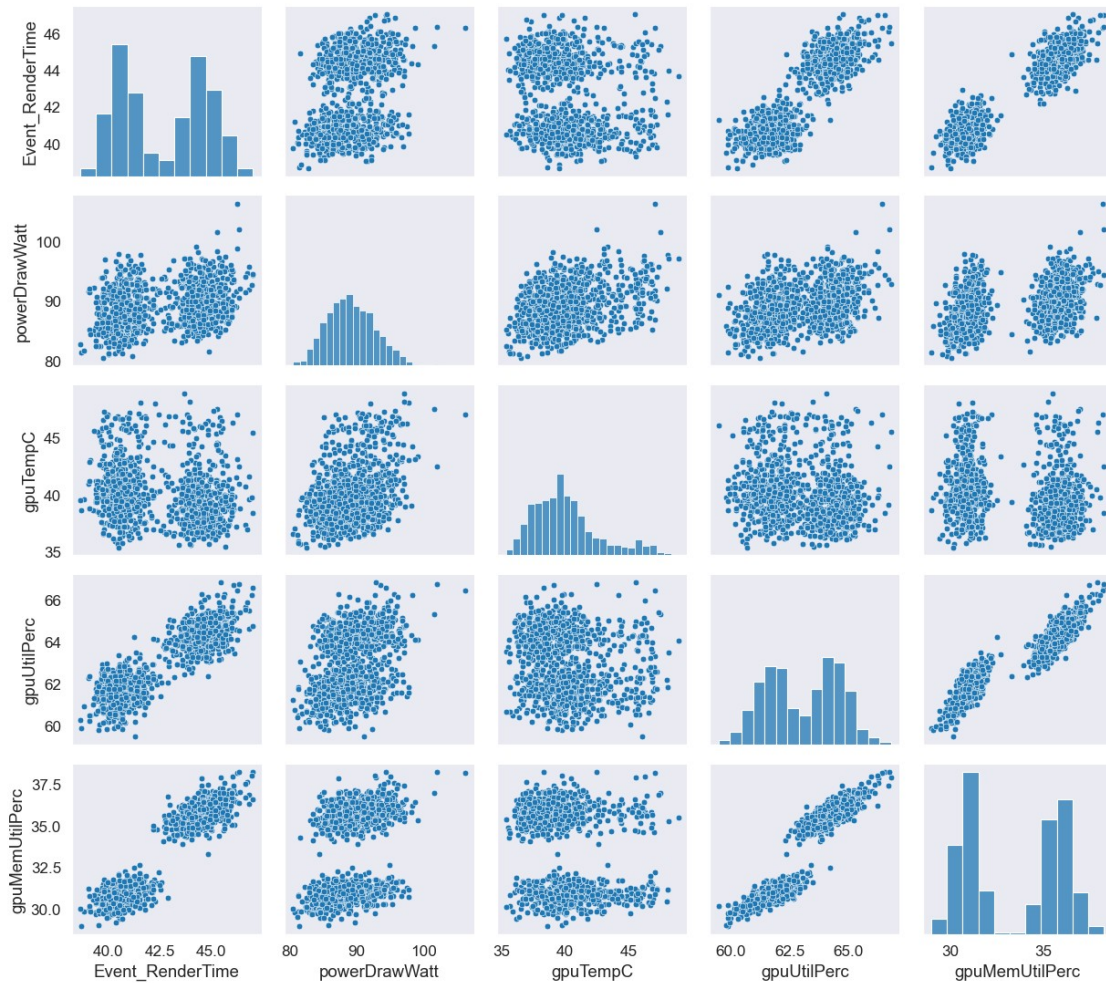
Plot a Joint Plot

```
sns.set_style('dark')
plt.figure(figsize=(7,6))
sns.set_context('paper', font_scale=1.4)
p = sns.pairplot(data=TP)
p.fig.tight_layout()
p.fig.subplots_adjust(top=0.87)
p.fig.suptitle("GPU Measure Comparison", y=1.0)
```

Text(0.5, 1.0, 'GPU Measure Comparison')

<Figure size 700x600 with 0 Axes>

GPU Measure Comparison



3. Identification of serial numbers of GPU with least performance

Worst performing GPU is calculated

```
Worst = """ select gpuSerial, Event_RenderTime As "Time" from TP
group by gpuSerial order by Time desc limit 10; """
```

```
GPU_Worst = pd.DataFrame(sqlfn(Worst))
```

```
GPU_Worst.index.name='Order'
```

```
GPU_Worst
```

	gpuSerial	Time
Order		
0	323617042821	47.038776
1	325017018645	47.013439
2	320218055378	46.993168
3	323217048793	46.853237
4	320118118728	46.833710
5	320218055639	46.743965
6	320118119009	46.675816
7	325117172889	46.660000
8	320118118802	46.399696
9	323617021137	46.350880

Create Fig and gridspec

```
import matplotlib.pyplot as plt
fg = plt.figure(figsize=(11,9), dpi= 85)
pg = plt.GridSpec(5, 4, hspace=0.6, wspace=0.3)
subp = fig.add_subplot(grid[:-2, :-1])
```

#Plotting Bar chart for displaying the time for each event

```
import seaborn as sns
g = sns.barplot(data=GPU_Worst, x="gpuSerial",
y="Time",order=GPU_Worst.sort_values('Time',ascending =
False).gpuSerial,color='yellow',alpha=.6)
```


Annotate Text

Annotate Text

```
for i, cty in enumerate(GPU_Worst.Time):
    g.text(i, cty+0.005, round(cty, 3),size = 11,
horizontalalignment='center')
```

Title, Label, Ticks and Ylim

```
g.set_title('GPU with Least Performance (Top 10)',
fontdict={'size':12})
g.set(ylabel='Run Time(Sec)',xlabel='GPU Serial Number',ylim=(46.3,
47.1),)
g.set_xticklabels(g.get_xticklabels(), rotation=45,
horizontalalignment='right')
# Add patches to color the X axis labels
plt.show()
```



TP

	hostname	Event_RenderTime \
0	04dc4e9647154250beeee51b866b0715000000	44.350327
1	04dc4e9647154250beeee51b866b0715000001	45.698051
2	04dc4e9647154250beeee51b866b0715000002	40.275421
3	04dc4e9647154250beeee51b866b0715000003	40.822983
4	04dc4e9647154250beeee51b866b0715000004	45.969166
...
1019	e7adc42d28814e518e9601ac2329c513000019	45.212002
1020	e7adc42d28814e518e9601ac2329c51300001A	40.272308
1021	e7adc42d28814e518e9601ac2329c51300001B	45.223820
1022	e7adc42d28814e518e9601ac2329c51300001C	44.953407
1023	e7adc42d28814e518e9601ac2329c51300001D	45.705097
	gpuUUID	gpuSerial
powerDrawWatt \		
0	GPU-a1119ee9-9cd1-919f-a479-b902142c717d	323217056165
95.868947		

```

1      GPU-3dc1601f-0e52-2e31-6b8d-8537d356b84d  323617042956
91.813693
2      GPU-e2deaca4-3041-9bbd-b6ac-6d36e09ab116  323617021222
82.537798
3      GPU-173a16e7-07e7-3677-b499-9e20353f81a3  323617021168
86.558581
4      GPU-c1bb7aec-1afc-3226-db53-58f077365bea  323217056664
94.292392
...
...
1019   GPU-53f2e984-36f3-a70f-0944-752868d6a833  320118119210
89.386562
1020   GPU-05e3f5b6-553b-1a37-bf7a-271c5b999a64  325117063265
82.943504
1021   GPU-8646ce3a-be99-02fa-a138-dafef7ea3cf2  325017048638
88.354997
1022   GPU-b0b09148-0762-390d-08d6-93c1a5ccd768  325117173230
88.411319
1023   GPU-0cee5a9f-749e-7780-791a-ff2b29590a38  320118119027
97.399313

```

```

      gpuTempC  gpuUtilPerc  gpuMemUtilPerc
0      43.525333      63.602667      35.876000
1      40.992000      64.630000      35.510000
2      38.048193      61.548862      30.505355
3      41.576282      60.127249      29.832778
4      39.742838      64.545636      35.549634
...
1019   39.095270      64.512325      36.449034
1020   37.712858      61.762825      30.990007
1021   38.569046      63.919947      35.593062
1022   37.205863      64.930047      36.289141
1023   39.816000      64.793333      36.607333

```

[1024 rows x 8 columns]

4. Interplay of temperature and power

```


# Plot a Joint Plot
sns.set_style('white')
plt.figure(figsize=(8,8))
sns.set_context('paper', font_scale=1.4)
p = sns.jointplot(x='powerDrawWatt', y='gpuTempC', data=TP,
kind='reg', joint_kws={'line_kws':{'color':'red'}})
p.fig.suptitle("Comparing GPU Temperature and power consumption")
p.ax_joint.set_xlabel('Power (in watt)')
p.ax_joint.set_ylabel('GPU Temperature (in C)')
p.fig.tight_layout()
p.fig.subplots_adjust(top=0.95)

```

```
print(f'The correlation coefficient between temperature and power =  
{np.round(np.corrcoef(TP.powerDrawWatt, TP.gpuTempC)[0,1], 2)}')
```

The correlation coefficient between temperature and power = 0.4

<Figure size 800x800 with 0 Axes>




5. variation in memory utilization percentage of the gpu with temperature

```
# Plot a Joint Plot
sns.set_style('white')
plt.figure(figsize=(8,8))
sns.set_context('paper', font_scale=1.4)
sns.scatterplot(x='gpuMemUtilPerc', y='gpuTempC', data=TP)
plt.title('Correlation coefficient between temperature and memory  
utilisation')
```



```
print(f'The correlation coefficient between temperature and memory  
utilisation = {np.round(np.corrcoef(TP.gpuMemUtilPerc, TP.gpuTempC)  
[0,1], 2)}')
```

The correlation coefficient between temperature and memory utilisation
= -0.18



Level_RT

	level	Event_RenderTime
0	4	52.181999
1	8	48.459820
2	12	42.581760


Plot barplot

```
sns.barplot(data=Level_RT, x="level", y="Event_RenderTime")
```



```
plt.title("Average Performance or Runtime of each Image Level",
size=20,fontweight='bold')
plt.xlabel("Image Levels", size = 14)
plt.ylabel("Runtime (in seconds)", size = 14)
plt.show()
```

Average Performance or Runtime of each Image Level



xy_level8

	x	y	Event_RenderTime
0	0	0	43.853001
1	0	1	24.809999
2	0	2	27.111000
3	0	3	35.481998
4	0	4	43.898998
...
251	15	11	64.322998
252	15	12	47.679001
253	15	13	47.452000
254	15	14	43.641998
255	15	15	45.969002

[256 rows x 3 columns]

```
# Plot seaborn heatmap
sns.set(font_scale=0.8)
sns.set_context('paper', font_scale=1.4)
fig, ax = plt.subplots(figsize=(15,10))
glue = xy_level8[["x","y","Event_RenderTime"]].pivot("x", "y",
"Event_RenderTime")
```

```
ax = sns.heatmap(glue, cmap='YlGnBu', annot=False, xticklabels=False,
yticklabels=False)
ax.tick_params()
```

```
# Change characteristics of the visual
plt.title("Heatmap based on Runtime for each co-ordinate of the
Rendered Tile", size=20,fontweight='bold' )
```

```
Text(0.5, 1.0, 'Heatmap based on Runtime for each co-ordinate of the
Rendered Tile')
```

